

Binomial logit t simulation Stan implementation

Divan A. Burger, Sean van der Merwe, and Emmanuel Lesaffre

2025-05-16

Table of contents

Introduction	1
Fixed effects Stan model	2
Base sample generation	2
Posterior simulation	3
Evaluation	3
Stan model with random intercept only	4
Posterior simulation	6
Evaluation	6
Adding random slopes	7
Evaluation	10

Introduction

This document provides introductory information and example code of how one might use Stan to build and fit a Binomial Generalized Linear Mixed Effects Model, but with additional Student-t innovations on the linear predictor.

```
options(scipen = 12)
library(tidyverse)
library(knitr)
```

Fixed effects Stan model

```
library(rstan)
mycores <- max(1, floor(parallel::detectCores(logical = FALSE)*0.8))
options(mc.cores = mycores)
rstan_options(auto_write = TRUE)
```

The model being used is a variation on the standard binomial regression model, but with a Bayesian implementation using objective priors and a Student-t distribution.

```
// This Stan block defines a binomial t regression model, by Sean van der Merwe, UFS
data {
  int<lower = 1> n_total;           // number of observations in total
  int<lower = 1> m;                 // binomial maximum parameter
  int<lower = 0, upper = m> y[n_total]; // observations
  int<lower = 1> n_var;             // number of explanatory variables, including intercept
  matrix[n_total, n_var] X;        // explanatory matrix, starting with intercept(1s)
}
// The parameters of the model
parameters {
  real<lower = 0> sigma;            // error scale parameter
  real<lower = 2> nu;               // error degrees of freedom
  vector[n_var] beta;              // coefficients
  vector[n_total] alpha;           // intermediate (nuisance) parameter
}
model {
  y ~ binomial_logit(m, alpha);     // likelihood
  alpha ~ student_t(nu, X*beta, sigma);
  target += log(nu) - 3*log(nu + 1) - log(sigma); // priors
}
```

Base sample generation

```
rbin_t <- function(n, m, mu, s, nu) {
  t_values <- rt(n, nu)*s + mu
  rbinom(n, m, plogis(t_values))
}

r_bin_t_data <- function(p) {
```

```

X <- cbind(matrix(1, p$n, 1), rnorm(p$n))
Z <- rnorm(p$n_sbj, 0, p$tau)
sbj <- sample(1:p$n_sbj, p$n, TRUE)
y <- rbin_t(p$n, p$m, X %*% p$beta + Z[sbj], p$s, p$nu)
list(n_total = p$n, m = p$m, y = y, n_var = length(p$beta), X = X, n_sbj = p$n_sbj, sbj = sbj)
}

```

```

params <- list(
  n = 200,
  m = 30,
  beta = c(1, 0.5),
  s = 1,
  nu = 6,
  tau = 0,
  n_sbj = 1
)
smpl <- r_bin_t_data(params)

```

Posterior simulation

```

fit1 <- bin_t_model |> sampling(smpl,
  pars = c('alpha'), include = FALSE,
  iter = 16000, chains = mycores)

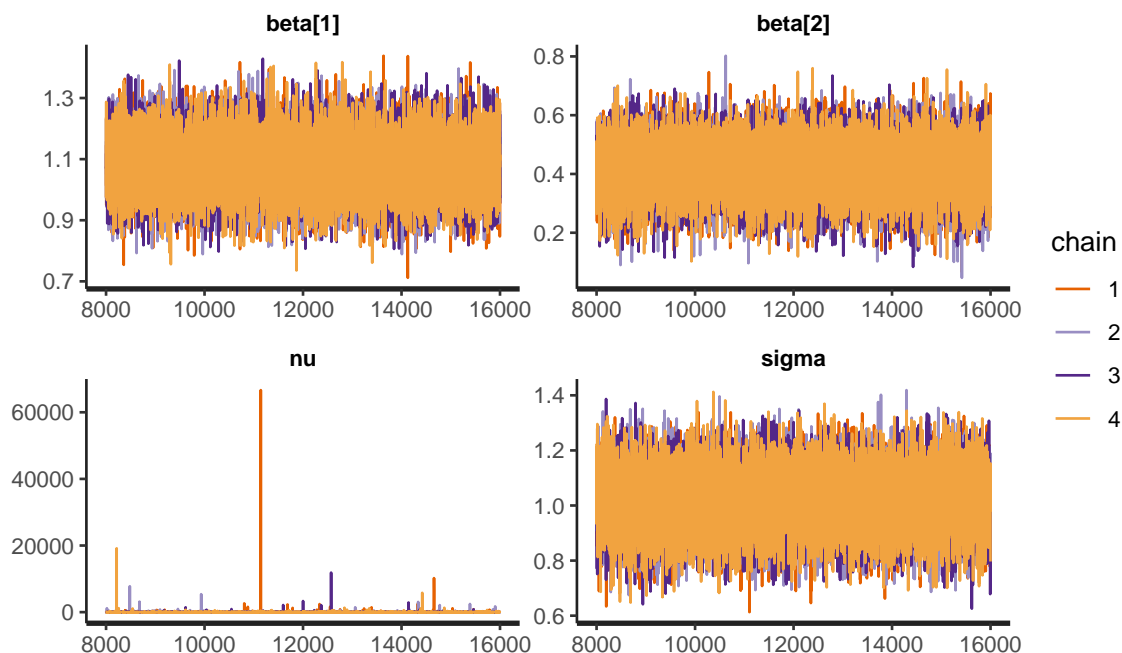
```

Evaluation

```

fit1 |> traceplot(pars = c('beta', 'nu', 'sigma'))

```



```
summary(fit1, pars = c('beta', 'nu', 'sigma'))$summary
```

	mean	se_mean	sd	2.5%	25%	50%
beta[1]	1.0837243	0.0004171409	0.08845227	0.9118969	1.0237860	1.0826017
beta[2]	0.4146311	0.0003937648	0.08508309	0.2463720	0.3576294	0.4148821
nu	21.2434332	2.6894982704	420.69732200	2.8279368	4.8078668	7.0211200
sigma	1.0126803	0.0008333114	0.10448535	0.8105439	0.9414000	1.0123514

	75%	97.5%	n_eff	Rhat
beta[1]	1.1435417	1.2576679	44962.71	0.9999565
beta[2]	0.4724247	0.5811322	46688.80	0.9999097
nu	11.6098155	66.7767855	24467.91	0.9999715
sigma	1.0836418	1.2167056	15721.58	1.0000329

Stan model with random intercept only

```
library(rstan)
mycores <- max(1, floor(parallel::detectCores(logical = FALSE)*0.8))
options(mc.cores = mycores)
rstan_options(auto_write = TRUE)
```

The model being used is a variation on the standard binomial regression model, but with a Bayesian implementation using objective priors and a Student-t distribution.

```
// This Stan block defines a binomial t regression model, by Sean van der Merwe, UFS
data {
  int<lower = 1> n_total;           // number of observations in total
  int<lower = 1> m;                 // binomial maximum parameter
  int<lower = 0, upper = m> y[n_total]; // observations
  int<lower = 1> n_var;             // number of explanatory variables, including intercept
  matrix[n_total, n_var] X;        // explanatory matrix, starting with intercept(1s)
  int<lower = 2> n_sbj;             // number of subjects
  int<lower = 1, upper = n_sbj> sbj[n_total]; // subject membership number
}
// The parameters of the model
parameters {
  real<lower = 0> sigma;            // error scale parameter
  real<lower = 2> nu;               // error degrees of freedom
  vector[n_var] beta;              // coefficients
  vector[n_total] alpha;           // intermediate (nuisance) parameter
  vector[n_sbj] sbj_int;           // subject intercept
  real<lower = 0> tau_sbj;          // between subject standard deviation
}
transformed parameters {
  vector[n_total] mu;
  mu = X*beta + sbj_int[sbj];
}
model {
  y ~ binomial_logit(m, alpha);     // likelihood
  alpha ~ student_t(nu, mu, sigma);
  sbj_int ~ normal(0, tau_sbj);
  target += log(nu) - 3*log(nu + 1) - log(sigma) - log(tau_sbj); // priors
}
```

```
params <- list(
  n = 200,
  m = 30,
  beta = c(1, 0.5),
  s = 1,
  nu = 6,
  tau = 0.5,
  n_sbj = 10
)
smpl <- r_bin_t_data(params)
```

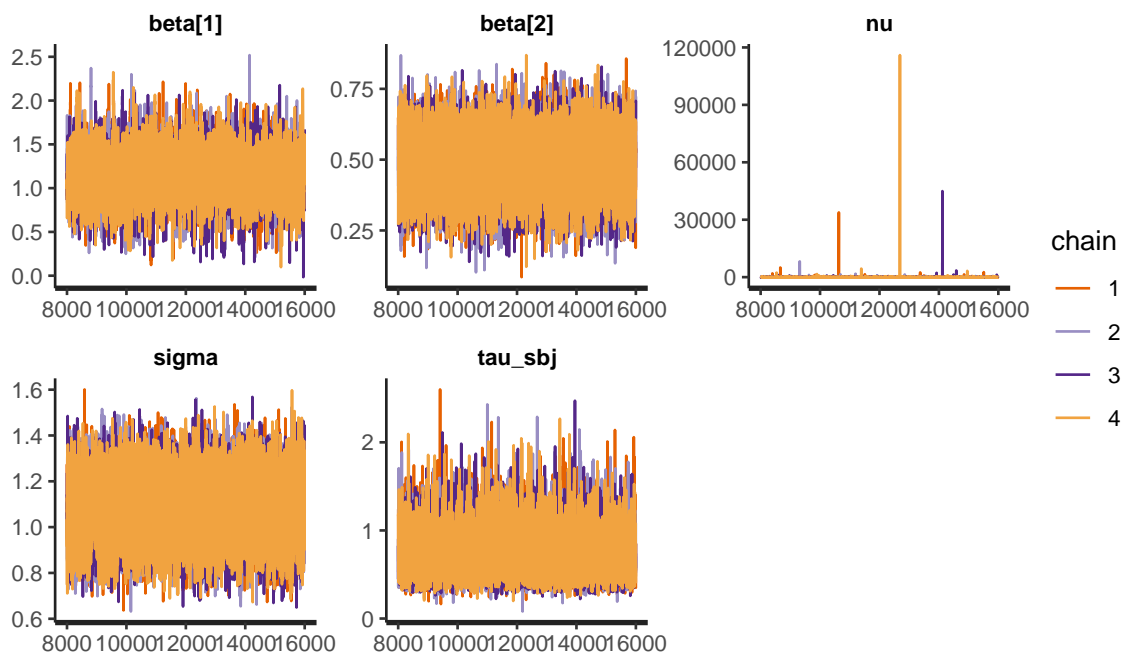
Posterior simulation

```
fit <- bin_t_model_rint |> sampling(smpl,  
  pars = c('alpha', 'mu'), include = FALSE,  
  iter = 16000, chains = mycores)
```

```
fit |> saveRDS("C:/temp/Simulations/bin_t_rint_stan_simulation.rds")
```

Evaluation

```
fit |> traceplot(pars = c('beta', 'nu', 'sigma', 'tau_sbj'))
```



```
summary(fit, pars = c('beta', 'nu', 'sigma', 'tau_sbj'))$summary
```

	mean	se_mean	sd	2.5%	25%	50%
beta[1]	1.1563283	0.0033044074	0.25408431	0.6426838	0.9994006	1.1551933
beta[2]	0.4795249	0.0005656486	0.09608752	0.2920067	0.4144875	0.4788924
nu	20.3874365	4.6968852213	741.15214832	2.2806730	3.5659077	5.0828092
sigma	1.0727852	0.0014779195	0.13251228	0.8228467	0.9798467	1.0704100
tau_sbj	0.7141516	0.0018555174	0.23057496	0.3764993	0.5545284	0.6766158

	75%	97.5%	n_eff	Rhat
beta[1]	1.3145870	1.6596277	5912.464	1.0008166
beta[2]	0.5443574	0.6677377	28856.293	1.0000394
nu	8.4196195	53.4823127	24899.742	0.9999957
sigma	1.1633928	1.3368114	8039.161	1.0003988
tau_sbj	0.8312491	1.2759265	15441.669	1.0000642

Adding random slopes

```
// This Stan block defines a binomial t regression model, adapted from
// code generated by the brms package
functions {
  /* compute correlated group-level effects
   * Args:
   *   z: matrix of unscaled group-level effects
   *   SD: vector of standard deviation parameters
   *   L: cholesky factor correlation matrix
   * Returns:
   *   matrix of scaled group-level effects
   */
  matrix scale_r_cor(matrix z, vector SD, matrix L) {
    // r is stored in another dimension order than z
    return transpose(diag_pre_multiply(SD, L) * z);
  }
}
data {
  int<lower=1> N; // total number of observations
  int<lower = 1> m; // binomial maximum parameter
  int<lower = 0, upper = m> Y[N]; // response variable
  int<lower=1> K; // number of population-level effects
  matrix[N, K] X; // population-level design matrix
  int<lower=1> Kc; // number of population-level effects after centering
  // data for group-level effects of ID 1
  int<lower=1> N_1; // number of grouping levels
}
```

```

    int<lower=1> M_1; // number of coefficients per level
    array[N] int<lower=1> J_1; // grouping indicator per observation
    // group-level predictor values
    vector[N] Z_1_1;
    vector[N] Z_1_2;
    int<lower=1> NC_1; // number of group-level correlations
}
transformed data {
    matrix[N, Kc] Xc; // centered version of X without an intercept
    vector[Kc] means_X; // column means of X before centering
    for (i in 2:K) {
        means_X[i - 1] = mean(X[, i]);
        Xc[, i - 1] = X[, i] - means_X[i - 1];
    }
}
parameters {
    vector[Kc] b; // regression coefficients
    real Intercept; // temporary intercept for centered predictors
    real<lower=0> sigma; // dispersion parameter
    vector<lower=0>[M_1] sd_1; // group-level standard deviations
    matrix[M_1, N_1] z_1; // standardized group-level effects
    cholesky_factor_corr[M_1] L_1; // cholesky factor of correlation matrix
    real<lower=1> nu; // degrees of freedom or shape
    vector[N] alpha; // intermediate (nuisance) parameter
}
transformed parameters {
    matrix[N_1, M_1] r_1; // actual group-level effects
    // using vectors speeds up indexing in loops
    vector[N_1] r_1_1;
    vector[N_1] r_1_2;
    real lprior = 0; // prior contributions to the log posterior
    // compute actual group-level effects
    r_1 = scale_r_cor(z_1, sd_1, L_1);
    r_1_1 = r_1[, 1];
    r_1_2 = r_1[, 2];
    lprior += student_t_lpdf(Intercept | 3, 23, 5.9);
    lprior += student_t_lpdf(sigma | 3, 0, 5.9)
        - 1 * student_t_lccdf(0 | 3, 0, 5.9);
    lprior += gamma_lpdf(nu | 2, 0.1)
        - 1 * gamma_lccdf(1 | 2, 0.1);
    lprior += student_t_lpdf(sd_1 | 3, 0, 5.9)
        - 2 * student_t_lccdf(0 | 3, 0, 5.9);
}

```



```

    lprior += lkj_corr_cholesky_lpdf(L_1 | 1);
  }
model {
  // initialize linear predictor term
  vector[N] mu = rep_vector(0.0, N);
  mu += Intercept + Xc * b;
  for (n in 1:N) {
    // add more terms to the linear predictor
    mu[n] += r_1_1[J_1[n]] * Z_1_1[n] + r_1_2[J_1[n]] * Z_1_2[n];
  }
  Y ~ binomial_logit(m, alpha);          // likelihood
  alpha ~ student_t(nu, mu, sigma);
  target += lprior;
  target += std_normal_lpdf(to_vector(z_1));
}
generated quantities {
  // actual population-level intercept
  real b_Intercept = Intercept - dot_product(means_X, b);
  // compute group-level correlations
  corr_matrix[M_1] Cor_1 = multiply_lower_tri_self_transpose(L_1);
  vector<lower=-1,upper=1>[NC_1] cor_1;
  // extract upper diagonal of correlation matrix
  for (k in 1:M_1) {
    for (j in 1:(k - 1)) {
      cor_1[choose(k - 1, 2) + j] = Cor_1[j, k];
    }
  }
}
}

```

```

stan_data <- list(N = smpl$n_total, m = smpl$m, Y = smpl$y, K = smpl$n_var,
                 X = smpl$X, Kc = smpl$n_var-1, N_1 = smpl$n_sbj, M_1 = smpl$n_var,
                 J_1 = smpl$sbj, Z_1_1 = smpl$X[,1], Z_1_2 = smpl$X[,2],
                 NC_1 = smpl$n_var - 1)

```


```

fit <- bin_t_model_rslope |> sampling(stan_data,
                                     pars = c('alpha', 'mu'), include = FALSE,
                                     iter = 16000, chains = mycores)

```

Evaluation

```
fit |> traceplot(pars = c('b', 'sigma', 'sd_1'))
```



```
summary(fit, pars = c('b', 'sigma', 'sd_1'))$summary
```

	mean	se_mean	sd	2.5%	25%	50%
b[1]	0.5099606	0.0009652586	0.1314722	0.25265999	0.4259583	0.5088327
sigma	1.1693370	0.0011346330	0.1137405	0.93038802	1.0980840	1.1729207
sd_1[1]	0.8392523	0.0029414558	0.2864831	0.43299844	0.6423364	0.7879782
sd_1[2]	0.2236565	0.0014262953	0.1572308	0.01102975	0.1049054	0.1990957
	75%	97.5%	n_eff	Rhat		
b[1]	0.5920380	0.7757086	18551.572	0.9999823		
sigma	1.2463073	1.3810868	10048.924	1.0002982		
sd_1[1]	0.9791324	1.5317442	9485.785	1.0003488		
sd_1[2]	0.3100315	0.5984530	12152.247	1.0001754		